

A Solver-Free Learning Framework for ILP

Rishabh Ranjan, Yatin Nandwani, Mausam, Parag Singla

BTP Final Presentation – 8 January 2022

Introduction

Motivation

Task

Perception \rightarrow Reasoning

Model

Neural \rightarrow Symbolic

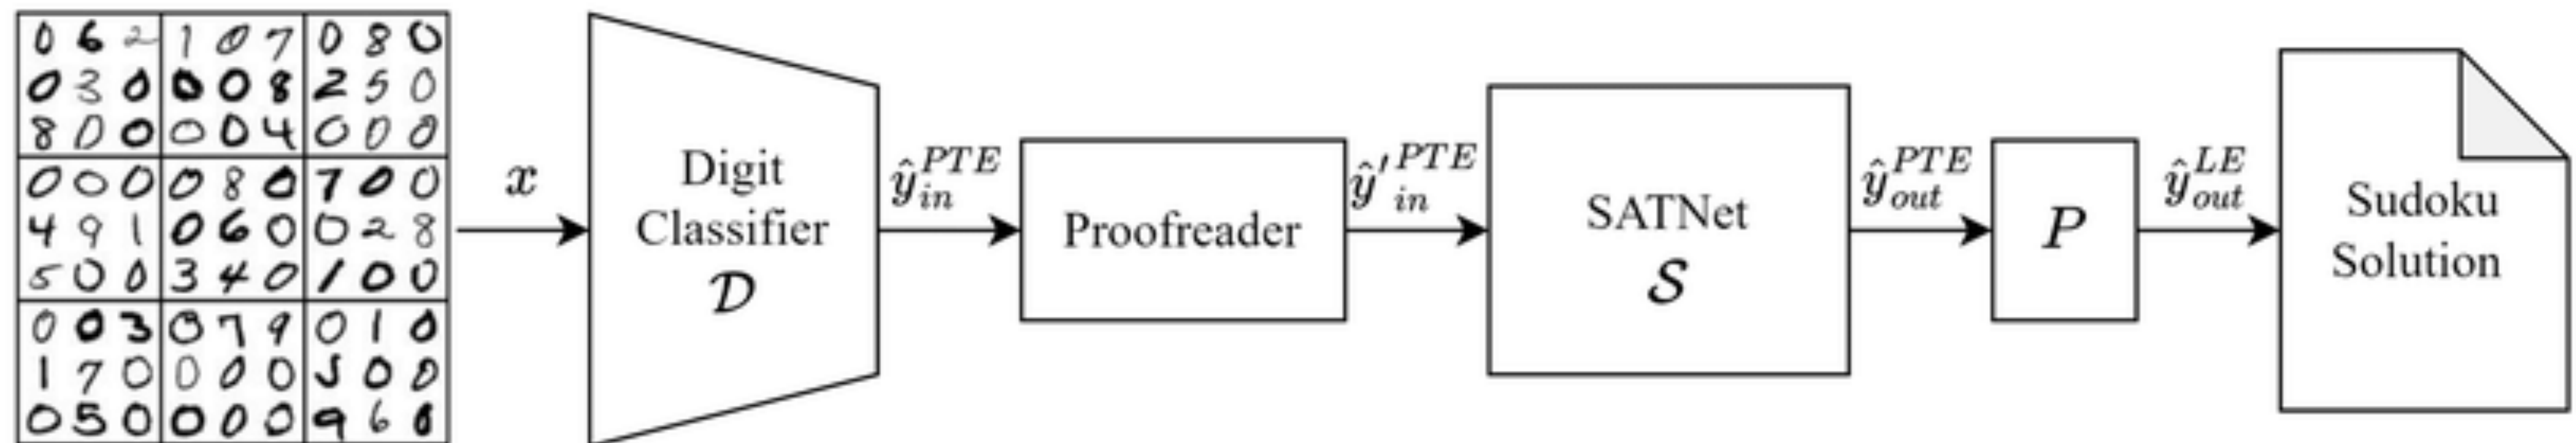


Figure from *Topan, Sever & Rolnick, David & Si, Xujie. (2021). Techniques for Symbol Grounding with SATNet.*

Integer Linear Programming

$x \rightarrow$ input

$A(x) \rightarrow$ constraint matrix

$b(x) \rightarrow$ bias vector

$c(x) \rightarrow$ cost vector

$y(A, b, c) \rightarrow$ solution vector / output

$n \rightarrow$ number of variables


$\min c \cdot y$ such that $Ay + b \geq 0, y \in \mathbb{Z}^n$

Related Work 1: CombOptNet

CombOptNet: Fit the Right NP-Hard Problem by Learning Integer Programming Constraints

Anselm Paulus, Michal Rolínek, Vít Musil, Brandon Amos, and Georg Martius

ICML 2021




Related Work 2: Constraint Mining


An Integer Linear Programming Framework for Mining Constraints from Data

Tao Meng, Kai-Wei Chang

ICML 2021



(a) Outer polytope (S_O)



(b) Inner polytope (S_I)

Related Work 3: Rectifier Networks

Learning Constraints for Structured Prediction Using Rectifier Networks

Xingyuan Pan, Maitrey Mehta, Vivek Srikumar

ACL 2020

- **Training:** 2-layer ReLU network
- **Inference:** network weights \rightarrow ILP constraints
- Arbitrary constraint parameterization **not supported**
- Learnable cost not **supported**

Methodology

Constraint Satisfaction Framework

Cost is ZERO


- Geometry & ML perspective:
constraint \leftrightarrow hyperplane \leftrightarrow linear classifier
all constraints \leftrightarrow polytope \leftrightarrow *veto-ensemble* of linear classifiers
- Supervision dichotomy:
Inference \implies combinatorial search via ILP solver
Training \implies binary classification on +ve & -ve examples
- Example source:
+ve example \leftarrow unique, target solution
-ve example(s) \leftarrow exponential, sampling strategy?

Proposed Loss Function


- $d_j(A, b, y) = \frac{A_j \cdot y + b_j}{\|A_j\|} \rightarrow$ signed distance from hyperplane
- $\ell(d_j, \pm) \rightarrow$ some binary classification loss
- $\mathcal{L}(A, b, y_+) = \max_j \ell(d_j, +)$
- $\mathcal{L}(A, b, y_-) = \min_j \ell(d_j, -)$
- Smooth max/min with temperature τ
- Hinge Loss with margin μ :
 $\ell(d_j, +) = \text{ReLU}(\mu - d_j), \quad \ell(d_j, -) = \text{ReLU}(\mu + d_j)$
- Cross-Entropy Loss:
 $\ell(d_j, +) = -\log(\sigma(d_j)), \quad \ell(d_j, -) = -\log(1 - \sigma(d_j))$

Solver-Free Negative Sampling


Nearest Neighbours



+ve & -ve examples




a polytope excluding nearest neighbours




a possible learnt polytope


Project and Round



an interim learnt polytope



non-integral projected negatives



a possible learnt polytope

Solver-Based Negative Sampling

- Solution to current **ILP** (if \neq target solution)
- Solution to **LP relaxation** (+ rounding, if \neq target solution)

Regularization via Variance Loss

- Constraint directions may collapse


- Minimize $\sum_i \sum_j \frac{A_i \cdot A_j}{\|A_i\| \|A_j\|}$

- Equivalently, minimize $\mathcal{L}_v = \left\| \sum_i \frac{A_i}{\|A_i\|} \right\|^2$

- Crucial for projected negative samples

Extending to Constraint Optimization

Capture cost via new constraint: $c \cdot y \leq c \cdot y_+$



Results

Experimental Setup

Task: Symbolic Sudoku

- **Sudoku sizes:** 4x4, 6x6
- **Sudoku modelling:** clues as constraints, clues as cost weights
- **Binary classification loss:** Cross-Entropy, Hinge
- **Negative sampling:** Nearest Neighbours, Projection, ILP solution, LP solution
- **Baselines:** CombOptNet, Rectifier Networks
- Test accuracy is either **100%** or **0%**

Takeaways 1: Learnability

- **CombOptNet:**
 - Ultimately **works** for all cases
 - Training with cost is visibly more **unstable**
- **Rectifier Networks:**
 - Clues as constraints: **not supported**
 - Clues as cost: loss decreases but fails to learn – **inadequate negatives**
- **Our approach:**
 - Cross-Entropy loss **doesn't work** except for some cases
 - All negative samplers **work** with Hinge Loss (with suitable hyperparams)

Takeaways 2: Training Time

- **Solver-based:**
 - training takes **hours** (~**11 hrs** for CombOptNet at 4x4 with cost), even with upto **64** cores
 - cost optimization is a lot **slower** than mere satisfaction (~**5x** slower for CombOptNet)
 - GPU acceleration is **pointless**
- **Solver-free:**
 - converge in **minutes** (~**3 mins** for our approach at 4x4), on a **single** core
 - **effective** GPU acceleration
 - cost optimization is as **fast** as satisfaction
 - required learning iterations are only **negligibly** worse

Next Steps

For BTP2

- Scale to 9x9 sudokus
- Perceptual component
- Other tasks

Thank You!
Questions?