

Learning CSPs via ILP solvers

Rishabh Ranjan, Yatin Nandwani, Mausam, Parag Singla

BTP Mid-Term Project Presentation

Motivation

- The learning aspect
- CSPs in the real world:
 - industrial scheduling
 - combinatorial optimization
 - program verification
 - ... and lots more
- Can neural networks learn CSPs from examples?

SUDOKU

To solve the puzzle, all the blank cells must be filled in using numbers from 1 to 9. Each number can appear once in each row, column and in the nine 3x3 blocks. You can successfully solve the puzzle just by using logic and the process of elimination.

Yesterday's Answer:
Solution to 2/10/11


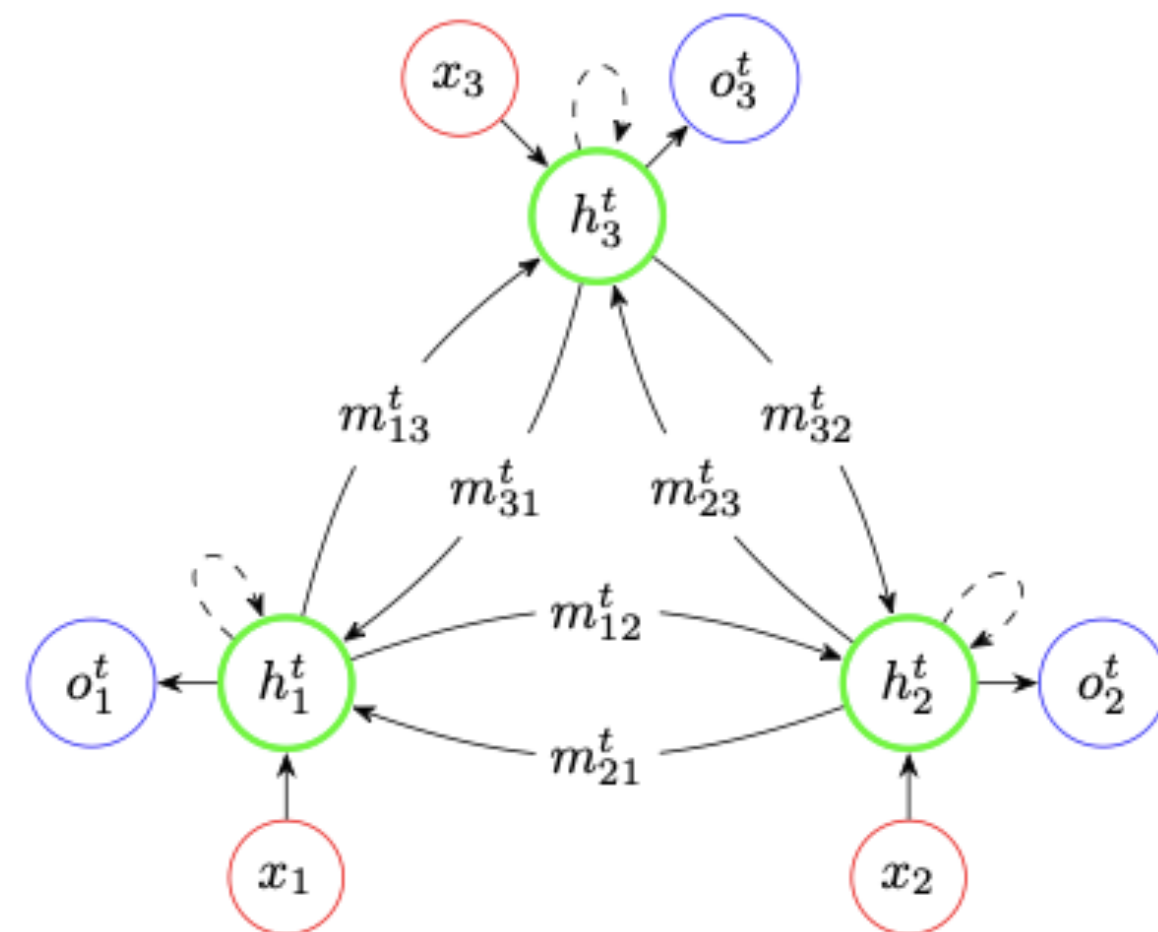
7	6	5	1	9	4	2	8	3
2	8	9	7	3	5	4	1	6
4	3	1	2	6	8	9	7	5
5	9	6	4	1	7	8	3	2
8	2	7	6	5	3	1	9	4
3	1	4	9	8	2	6	5	7
9	4	3	5	2	1	7	6	8
6	5	2	8	7	9	3	4	1
1	7	8	3	4	6	5	2	9

© 2011 Janric Enterprises Dist. by creators.com

2/11/11
SILVER

Approach 1: Neural Reasoner


Recurrent Relational Networks, Palm et. al., NeurIPS 2018



- Learns to reason from scratch
- Problem learning and solution learning are coupled together
- No match for decades old symbolic approaches

Approach 2: Symbolic Reasoner + RL


End-to-End Neuro-Symbolic Architecture for Image-to-Image Reasoning Tasks,
Agarwal et. al., unpublished, 2021



- Difficult to train: sparse rewards
- Problem specific solver: intractable action space for general-purpose solver

Approach 3: Symbolic Reasoner + Backprop

CombOptNet: Fit the Right NP-Hard Problem by Learning Integer Programming Constraints,
Paulus et. al., ICML 2021



- Backprop via informative gradients for black-box ILP solvers
- Any problem in NP is reducible to ILP (NP-completeness)
- Industry-grade solvers available: Gurobi

Limitations of Demonstrated Tasks


- Naive parametrisation of constraint matrix
- **Synthetic:** random dense constraints
- **Small:** max 8 constraints over 16 variables
- **Input dependence:**
 - only 1 setting, with only 1 constraint
 - static constraints in most tasks
- **Instance-size dependence**
- **Slow training:** 6 hrs for 4000 weight updates on one task (prior art took 1 hr)

Real-World CSPs

(Toy CSP but representative)

- ILP for 4x4 Sudoku: 82 constraints over 64 variables
- Input dependent constraints
- **Empirical Results:**
 - 4x4, static constraints only:
 - 100% accuracy
 - 4x4, fully learnable / 9x9, static constraints only:
 - trivial baseline accuracy


Accuracy Curves



Smoothed Digit Accuracy vs Training Steps

(BLUE: 4x4 sudoku, fully learnable,


ORANGE: 9x9 sudoku, only static constraints learnable):
stuck as trivial baseline accuracies of 25% and 40%



Smoothed Board Accuracy vs Training Steps
(4x4 sudoku, only static constraints learnable):
100% generalisation accuracy

Technical Challenge 1: Solver Interaction

- Forward pass calls ILP solver
- Randomly learnt constraints at training
- Unaligned to solver heuristics
- Explosive solving times
- Extremely slow training
- GPU acceleration is useless



Digit Accuracy vs Training Steps
(4x4 sudoku, fully learnable):
Observe how training steps slow down.

Tricks to Speed Up Training

- **Initialisation** with trivially feasible constraints
- Adaptive **solver timeouts**: low timeout initially, increase gradually
- **Solution hinting**:
 - solution of same training instance from last epoch
 - gold solution at later stages of training
 - advanced **mini-batching**: consecutive batches can have same instances
- Priors can help generate easier constraints

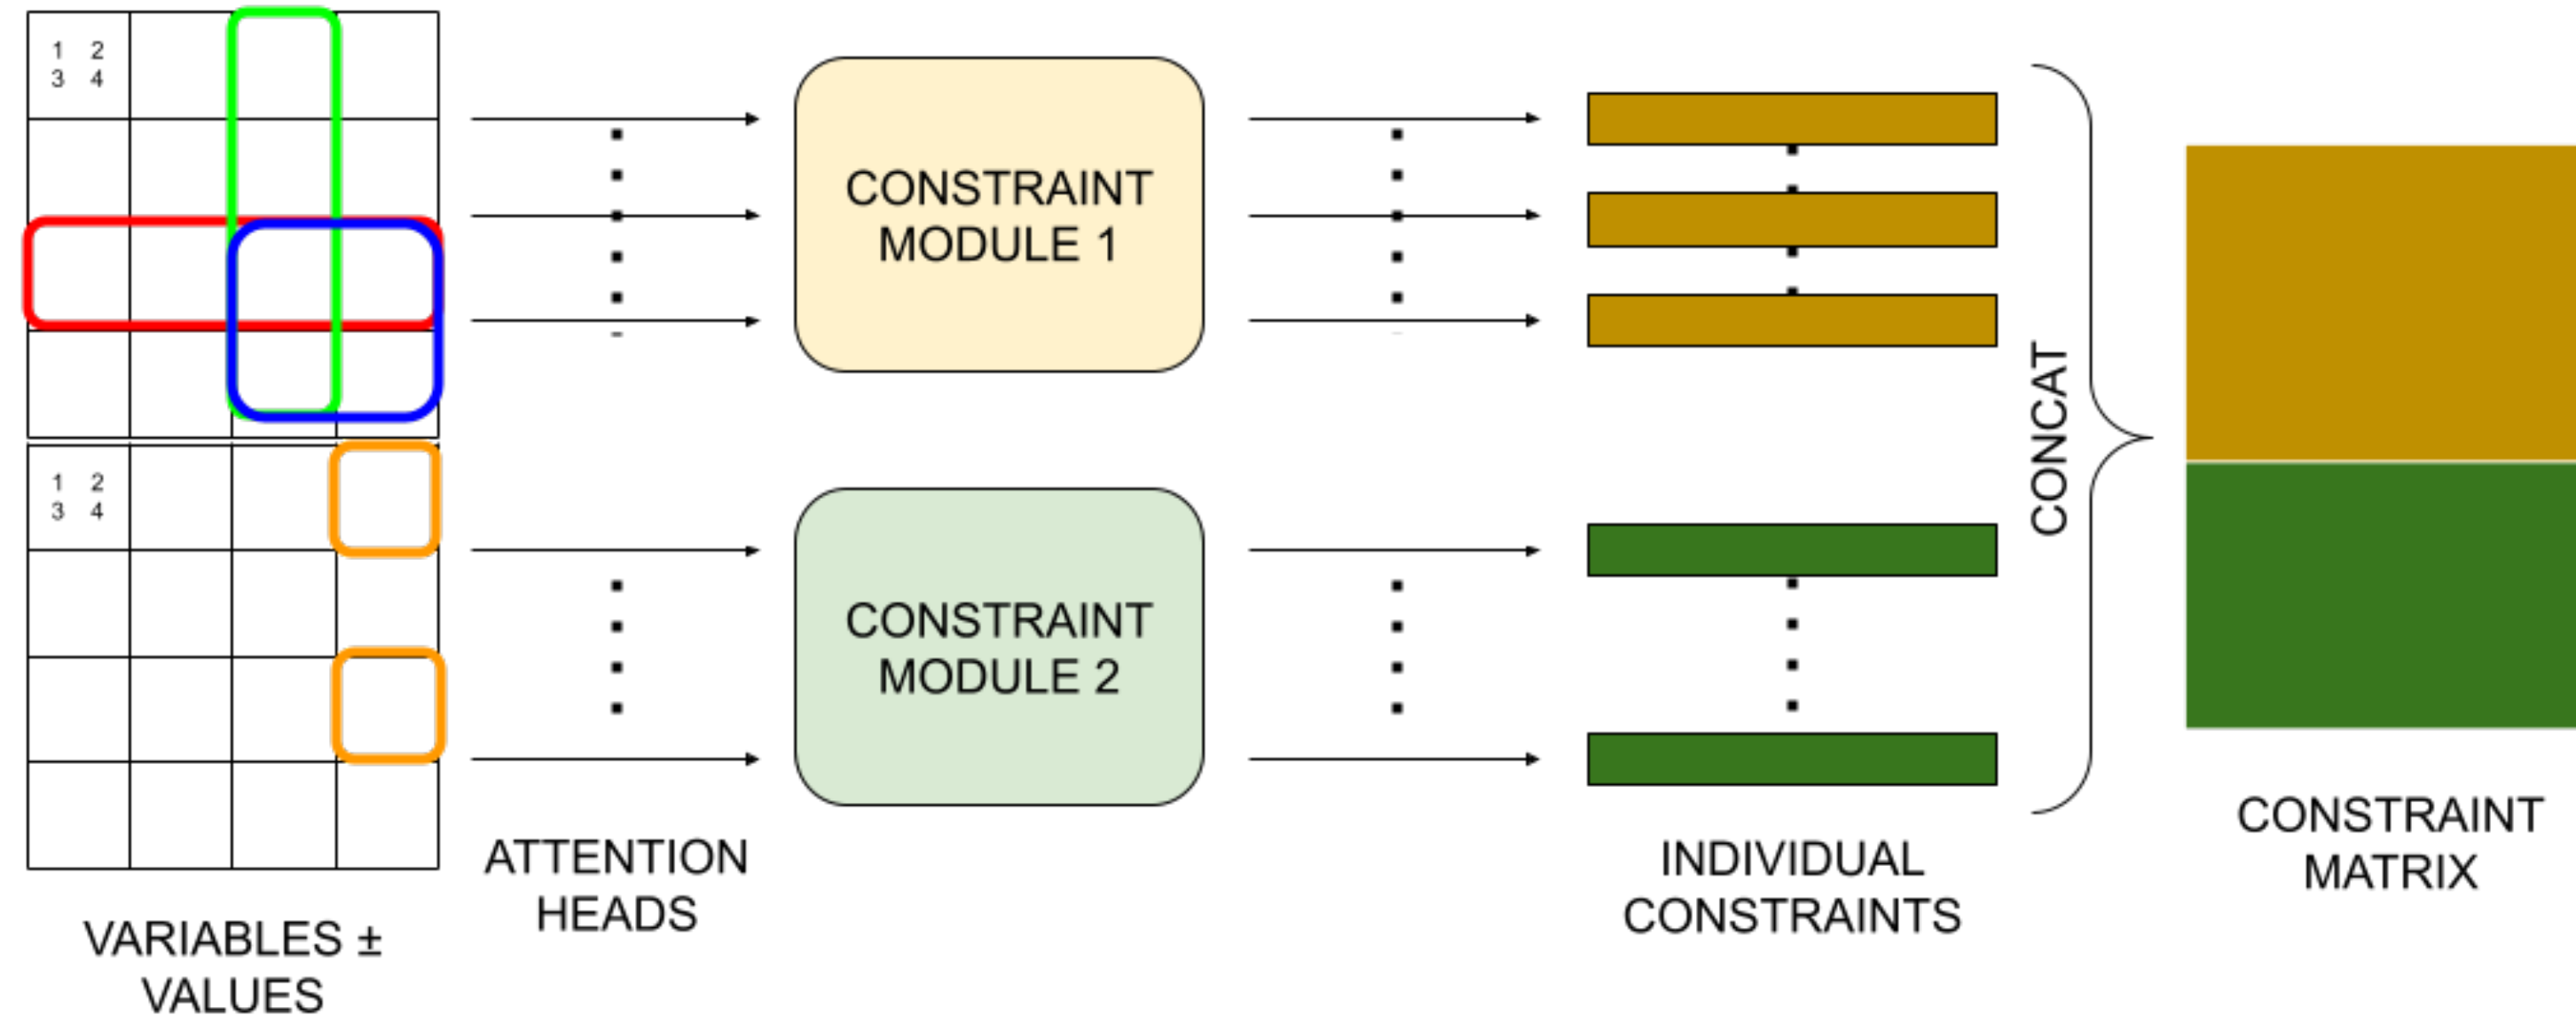
Technical Challenge 2: Parameterisation

- Each coefficient is learnt independently
- Example 1:
 - **ONE-HOT** constraint: $x + y + z = 1$ ($x, y, z \in \{0, 1\}$)
 - ILP form: (1) $x + y + z \leq 1$, (2) $(-x) + (-y) + (-z) \leq -1$
- Example 2:
 - **ALL-DIFF** constraint:
 - 9 rows, 9 columns, 9 boxes

The Real-World Redemption

- Priors:
 - **Sparsity:** ~99% for sudoku
 - **Lifting:** instantiations of the same constraints over variable permutations
 - **High-level constraints:**
 - Google OR Tools: `AddAllDifferent`, `AddReservoirConstraint`, `AddAllowedAssignments`, `AddBoolOr`, `AddMaxEquality`
 - Microsoft Z3 SMT Solver: `AtLeast(k)`, `AtMost(k)`
 - **Compositionality:** high-level from low-level constraints
- Heavily engineered solvers: Heuristics and optimisations for real-world instances


Proposed Architecture for Parameterisation



- **Attention** over variables (and values for input-dependence) as taking **soft subsets**
- **Learnable** or **hand-crafted** constraint modules
- **Permutation invariance:**
 - limitation in general
 - practically useful as high-level constraints tend to be permutation invariant

Alternative Solutions / Baselines / Directions

- **L1 regularisation:** sparsity prior
- **Deconvolution networks:** weight sharing in generation
- Restore generality using sequence selector modules
- Add a **compositionality** prior by going deeper



A deconvolution filter

Desired / Expected Benefits

- **Data efficiency:** learn only to formulate the problem
- **Faster forward pass:** priors aligned with ILP heuristics
- **Strong generalisation:** train on easy instances, test on hard instances
- **Domain knowledge:** enforce hard constraints directly
- **Instance size generalisation:** train on small instances, test on larger instances

The Plan Ahead

- **DONE:**
 - Problem formulation
 - Major implementation
 - Baseline Experiments
- **TODO:**
 - Thorough experiments on proposed architecture
 - Iterative development of more effective ideas

Thank You!

Questions?